

RiverCompute Intelligence — Build Workspace

Core Product Thesis

RiverCompute Intelligence is a planning intelligence platform for hyperscale AI infrastructure, focused on the real constraint behind river-adjacent compute expansion: heat rejection capacity, seasonal thermal stress, infrastructure throughput, and cumulative watershed risk.

Instead of treating water availability as a simple volume problem, the platform evaluates whether a river segment can safely support data center buildout without triggering ecological harm, permitting delays, or regulatory stress.

Primary Users

- Hyperscale data center developers and operators
- Utilities and municipal water authorities
- State governments and environmental regulators
- River basin authorities
- AI infrastructure investors
- Sustainability and policy research teams

High-Value MVP Positioning

RiverCompute Intelligence helps developers, regulators, utilities, and investors answer one question:

How much AI compute can a river segment safely support before heat, water withdrawal, seasonal stress, or infrastructure limits create unacceptable ecological or regulatory risk?

MVP Feature Stack

1. Watershed Compute Zoning Dashboard

Classifies river segments into compute suitability zones.

Inputs

- River flow
- Seasonal variability
- Facility count
- Industrial density
- Thermal sensitivity
- Existing infrastructure capacity

Outputs

- Zone A: Safe / preferred
- Zone B: managed caution
- Zone C: regulatory stress
- Zone D: overload / unsuitable

2. Thermal Load Risk Scoring

Estimates thermal risk from cooling load and seasonal river stress.

Inputs

- Facility megawatts
- Cooling architecture
- River baseline temperature
- Seasonal condition
- Discharge limits

Outputs

- Thermal risk score
- Risk interpretation
- Mitigation requirement

3. Cooling Architecture Comparator

Compares once-through, evaporative, hybrid, and dry cooling strategies.

Outputs

- Cooling risk profile
- Water tradeoff
- Energy tradeoff
- Recommended cooling option

4. Seasonal Stress Simulator

Models low-flow and high-temperature conditions, especially summer risk windows.

Outputs

- Summer stress multiplier
- Winter relief factor
- Seasonal permitting risk

5. Water-to-Query Calculator

Translates AI query volume and energy demand into estimated water impact.

Inputs

- Query count
- Wh/query
- Data center WUE
- Cooling type

Outputs

- Liters/gallons water estimate
- Query-to-water comparison
- Scenario summary

6. MSHDI Calculator

Maximum Sustainable Hyperscale Density Index.

Purpose Scores the sustainable density of compute infrastructure along a river segment.

Outputs

- MSHDI score
- Sustainable / Caution / Stress / Overload category
- Buildout recommendation

7. Regulatory Readiness Report

Generates a policy-ready report for siting, permitting, and environmental review.

Outputs

- Risk summary
- Clean Water Act readiness notes
- Mitigation recommendations
- Exportable report preview

8. Mitigation Planner

Recommends controls based on zone, risk score, cooling type, and watershed condition.

Outputs

- Cooling strategy adjustments
- Monitoring checklist

- Seasonal operating constraints
- Regulator-facing mitigation language

Decision Intelligence Model

The platform can use a weighted scoring model:

- Thermal Load: 35%
- Water Withdrawal Density: 25%
- Seasonal Overlap Risk: 20%
- Infrastructure Redundancy: 20%

Risk Bands

Score	Classification	Meaning
0–40	Safe	Strong buildout candidate
40–70	Managed Risk	Buildout possible with mitigation
70–100	Regulatory Stress	High review/permitting burden
100+	Overload	Unsuitable without major redesign

Knowledge Base Architecture

Core Tables / Collections

River Segments

- segment_id
- river_name
- state
- county
- latitude
- longitude
- mean_flow_cfs
- low_flow_cfs
- baseline_temp_c
- summer_temp_c
- dissolved_oxygen_mg_l
- ecological_sensitivity_score

Data Center Facilities

- facility_id
- operator
- status
- location

- megawatts
- cooling_type
- water_source
- estimated_wue
- discharge_method
- permit_status

Cooling Architectures

- cooling_type
- water_intensity
- thermal_discharge_risk
- energy_penalty
- seasonal_sensitivity
- best_use_case

Regulatory Thresholds

- jurisdiction
- discharge_temp_limit
- max_delta_t_c
- withdrawal_limit
- reporting_requirement
- permitting_notes

Scenario Runs

- scenario_id
- segment_id
- proposed_facility_mw
- cooling_type
- season
- risk_score
- mshdi_score
- recommendation
- report_output

AI Agent Design

Agent Name Ideas

- RiverIntel Agent
- HydroCompute Analyst
- Watershed Compute Advisor
- Thermal Risk Copilot
- BasinBuild Intelligence Agent

Agent Capabilities

The agent should be able to:

1. Pull real-time or recent public hydrological data.
2. Explain why a river segment is classified as safe, caution, stress, or overload.
3. Compare cooling architectures for a proposed facility.
4. Generate a regulator-facing readiness summary.
5. Save scenario runs into a project journal or case file.
6. Produce investor-facing risk language.
7. Flag missing data before generating conclusions.
8. Convert raw data into structured datasets for dashboard use.

Agent Prompt Frame

You are RiverCompute Intelligence Agent, a watershed-based AI infrastructure planning assistant. Your job is to help developers, regulators, utilities, and investors evaluate hyperscale data center siting risk using hydrological, thermal, ecological, regulatory, and infrastructure data. You must separate known data from assumptions, flag uncertainty, and generate clear risk classifications with practical mitigation recommendations.

Suggested MVP Workflow

1. User selects or searches a river segment.
2. User enters proposed facility load in MW.
3. User selects cooling architecture.
4. System pulls or references hydrological and regulatory data.
5. System calculates:
6. Thermal Load Risk
7. Water Withdrawal Density
8. Seasonal Overlap Risk
9. Infrastructure Redundancy
10. MSHDI score
11. Dashboard classifies the segment.
12. AI agent generates:
13. Plain-language explanation
14. Technical report summary
15. Mitigation checklist
16. Policy-ready export
17. User saves the run to the project journal / knowledge base.

90-Day Build Plan

Days 1–30: Data + Model Foundation

- Define schema for river segments, facilities, cooling types, and regulatory thresholds.
- Build sample dataset for Ohio River / Louisville case study.
- Create initial Thermal Load Risk formula.
- Create initial Water-to-Query formula.
- Create MSHDI scoring logic.
- Define risk band language.

Days 31–60: Dashboard + Report Generator

- Build Watershed Compute Zoning Dashboard.
- Add scenario input form.
- Add cooling architecture comparison view.
- Add seasonal stress simulator.
- Build regulatory readiness report generator.
- Add save/export function.

Days 61–90: Agent + Pilot Package

- Add AI research agent with source-grounded retrieval.
- Add case journal / scenario history.
- Package Louisville case study.
- Create investor/regulator demo deck.
- Prepare pilot outreach list.
- Test with sample developer, utility, and policy workflows.

Immediate Next Build Deliverables

1. **Structured dataset templates** for river segments, facilities, cooling types, regulations, and scenario runs.
 2. **Knowledge base schema** for retrieval-augmented AI responses.
 3. **Risk scoring formulas** for Thermal Load Risk and MSHDI.
 4. **Agent system prompt** and action map.
 5. **Dashboard wireframe** for the MVP.
 6. **Regulatory Readiness Report template.**
 7. **Pilot case study outline** for Louisville / Ohio River.
-

MVP Build Spec — RiverCompute Intelligence

MVP Goal

Build a functional prototype that proves RiverCompute Intelligence can evaluate whether a river segment can safely support hyperscale AI infrastructure by combining hydrological conditions, cooling architecture, thermal discharge risk, infrastructure capacity, and regulatory readiness into one decision-support workflow.

MVP User Story

A developer, regulator, utility planner, or investor selects a river segment, enters a proposed data center load, chooses a cooling architecture, and receives a clear risk classification, buildout recommendation, mitigation checklist, and exportable regulatory-readiness summary.

MVP Core Workflow

1. Select river segment.
2. Enter proposed facility size in megawatts.
3. Choose cooling architecture.
4. Select season or stress condition.
5. Run risk model.
6. View classification:
 7. Safe
 8. Managed Risk
 9. Regulatory Stress
 10. Overload
11. Generate mitigation plan.
12. Export a regulatory-readiness report.
13. Save scenario to project journal / case history.

MVP Pages / Screens

1. Landing / Project Dashboard

Purpose: Give the user a fast overview of active river-segment scenarios.

Includes:

- Project name
- Total scenarios
- Safe zones count

- Managed risk count
- Regulatory stress count
- Overload count
- Recent scenario runs
- “New Scenario” button

2. River Segment Explorer

Purpose: Let users review river segment conditions before running a scenario.

Includes:

- River segment list
- Segment map placeholder
- Mean flow
- Low-flow condition
- Baseline temperature
- Summer temperature
- Ecological sensitivity score
- Existing industrial/data center density
- Regulatory jurisdiction

3. Scenario Builder

Purpose: Main input screen for proposed data center buildout.

Inputs:

- River segment
- Proposed facility MW
- Cooling type
- Season
- Facility count in segment
- Estimated WUE
- Infrastructure redundancy level
- Mitigation controls selected

Primary button:

- **Run RiverCompute Analysis**

4. Risk Intelligence Results

Purpose: Show calculated scores and explain them.

Outputs:

- Thermal Load Risk score

- Water Withdrawal Density score
- Seasonal Overlap Risk score
- Infrastructure Redundancy score
- Composite risk score
- MSHDI score
- Classification band
- Recommendation
- Confidence level
- Missing data warnings

5. Cooling Architecture Comparator

Purpose: Compare cooling tradeoffs.

Cooling options:

- Once-through cooling
- Evaporative cooling
- Hybrid cooling
- Dry cooling

Outputs:

- Water intensity
- Heat discharge risk
- Energy penalty
- Best-fit conditions
- Recommended strategy

6. Mitigation Planner

Purpose: Translate risk into action.

Outputs:

- Required monitoring
- Seasonal operating limits
- Cooling modifications
- Thermal discharge controls
- Regulator-facing mitigation language
- Developer action checklist

7. Regulatory Readiness Report

Purpose: Generate a formal planning summary.

Report sections:

- Project overview
- River segment profile
- Proposed compute load
- Cooling architecture
- Thermal risk summary
- Seasonal stress summary
- MSHDI interpretation
- Mitigation plan
- Regulatory readiness status
- Decision recommendation

8. Scenario Journal

Purpose: Save analysis runs for future research, policy review, or investor documentation.

Includes:

- Scenario name
- Date created
- River segment
- MW load
- Risk score
- MSHDI score
- Recommendation
- Notes
- Export button

MVP Dataset Templates

river_segments.csv

Field	Type	Description
segment_id	string	Unique river segment ID
river_name	string	River name
state	string	State
county	string	County
latitude	number	Segment latitude
longitude	number	Segment longitude
mean_flow_cfs	number	Average river flow
low_flow_cfs	number	Low-flow condition

Field	Type	Description
baseline_temp_c	number	Normal water temperature
summer_temp_c	number	Summer water temperature
dissolved_oxygen_mg_l	number	Dissolved oxygen level
ecological_sensitivity_score	number	0–100 sensitivity score
infrastructure_capacity_score	number	0–100 capacity score
regulatory_jurisdiction	string	Agency or jurisdiction

facilities.csv

Field	Type	Description
facility_id	string	Facility ID
facility_name	string	Facility name
operator	string	Operator or owner
segment_id	string	Linked river segment
status	string	Proposed, active, planned, retired
megawatts	number	Compute load
cooling_type	string	Cooling architecture
estimated_wue	number	Water usage effectiveness
discharge_method	string	Discharge approach
permit_status	string	Permit stage

cooling_architectures.csv

Field	Type	Description
cooling_type	string	Cooling architecture
water_intensity_score	number	0–100 water intensity
thermal_discharge_score	number	0–100 heat discharge risk
energy_penalty_score	number	0–100 energy tradeoff
seasonal_sensitivity_score	number	0–100 seasonal vulnerability
best_use_case	string	Best-fit scenario

regulatory_thresholds.csv

Field	Type	Description
jurisdiction	string	Agency / region
max_delta_t_c	number	Maximum allowable river temperature increase
discharge_temp_limit_c	number	Discharge temperature limit
withdrawal_limit_mgd	number	Water withdrawal threshold
reporting_requirement	string	Reporting requirement
permitting_notes	text	Notes for report generation

scenario_runs.csv

Field	Type	Description
scenario_id	string	Scenario ID
date_created	date	Creation date
segment_id	string	River segment
proposed_mw	number	Proposed compute load
cooling_type	string	Cooling architecture
season	string	Seasonal condition
thermal_load_score	number	Thermal score
water_withdrawal_score	number	Water score
seasonal_overlap_score	number	Seasonal score
infrastructure_redundancy_score	number	Infrastructure score
composite_risk_score	number	Final score
mshdi_score	number	Density index
classification	string	Safe / Managed Risk / Regulatory Stress / Overload
recommendation	text	Decision recommendation

MVP Scoring Logic

Composite Risk Score

Composite Risk Score =

- Thermal Load Risk $\times 0.35$
- Water Withdrawal Density $\times 0.25$
- Seasonal Overlap Risk $\times 0.20$
- Infrastructure Redundancy Risk $\times 0.20$

Classification Bands

Score	Classification	Recommendation
0–40	Safe	Buildout candidate
40–70	Managed Risk	Build with mitigation
70–100	Regulatory Stress	High regulatory burden
100+	Overload	Redesign or avoid

Button Trigger Actions

New Scenario

Opens the Scenario Builder.

Run RiverCompute Analysis

Runs the scoring model and generates the risk result.

Compare Cooling Options

Runs cooling architecture comparison for the selected facility and river segment.

Generate Mitigation Plan

Creates checklist based on risk classification and cooling type.

Generate Regulatory Report

Creates a structured report from the scenario data.

Save to Scenario Journal

Stores the scenario run in the project history.

Export Report

Exports the regulatory-readiness report as PDF, DOCX, or Markdown.

MVP AI Agent Actions

Research River Segment

Pulls or retrieves available public data about a river segment.

Explain Risk Score

Explains what drove the risk classification.

Suggest Cooling Strategy

Recommends cooling architecture based on water, heat, and energy tradeoffs.

Draft Regulatory Summary

Generates formal language for regulators or permitting review.

Identify Missing Data

Flags where assumptions were used or where real data is needed.

Save Research Note

Adds findings to the Scenario Journal.

MVP Success Criteria

The MVP is successful if it can:

1. Run a complete scenario from input to recommendation.
2. Produce a credible composite risk score.
3. Classify a river segment into a clear risk band.
4. Explain why the classification was reached.
5. Compare cooling architectures.
6. Generate a mitigation checklist.
7. Export a regulatory-readiness report.
8. Save scenario runs for future review.

First Demo Case Study

Use the Ohio River / Louisville region as the first pilot case because it aligns with the platform's current strategic framing and provides a strong demonstration of watershed-based compute governance.

Build Priority

Must Build First

- Scenario Builder
- Risk scoring model
- River segment dataset
- Results dashboard
- Regulatory report generator
- Scenario journal

Build Second

- Live data integrations
- Map-based zoning
- AI research agent
- Cooling comparator
- Investor risk view

Build Later

- Multi-watershed national map
 - Real-time monitoring feeds
 - API for regulators/developers
 - Enterprise permissioning
 - Automated policy brief generation
-

Sustainability Relevance Layer

Why Sustainability Is Central to RiverCompute Intelligence

RiverCompute Intelligence should treat sustainability as a core decision layer, not as a marketing add-on. Hyperscale AI infrastructure creates overlapping sustainability pressures across electricity demand, water consumption, heat rejection, aquatic ecosystem protection, land-use planning, grid reliability, and environmental permitting.

The platform's strongest sustainability relevance is its ability to translate compute growth into watershed-level impact. Instead of asking only whether a region has enough water, RiverCompute asks whether the local ecosystem, infrastructure, and regulatory environment can absorb the heat, withdrawal, discharge, and seasonal stress created by concentrated AI compute.

Sustainability Problem Frame

AI infrastructure sustainability is not just about carbon emissions. It includes:

- Electricity demand growth
- Renewable energy availability
- Water withdrawal
- Consumptive water use
- Thermal discharge
- Wastewater quality
- River temperature stress
- Dissolved oxygen risk
- Aquatic habitat protection
- Seasonal drought and low-flow vulnerability
- Grid congestion
- Community and environmental justice concerns
- Permit-readiness and regulatory transparency

Sustainability Value Proposition

RiverCompute Intelligence helps stakeholders plan AI infrastructure that is:

1. **Ecologically safer** — by identifying where compute load may create thermal or aquatic stress.
2. **Regulator-ready** — by translating technical risk into permitting language and mitigation plans.
3. **Water-aware** — by separating direct water use, indirect energy-related water use, and thermal discharge impacts.
4. **Energy-conscious** — by showing tradeoffs between cooling efficiency, water demand, and power penalty.
5. **Climate-resilient** — by modeling seasonal low-flow and high-temperature stress scenarios.
6. **Infrastructure-aligned** — by preventing overconcentration of compute facilities along river segments that lack sufficient capacity.

Sustainability Metrics to Add to MVP

Water Metrics

- Site Water Usage Effectiveness
- Estimated consumptive water use
- Withdrawal volume
- Discharge volume
- Water source type
- Reclaimed vs potable water percentage
- Low-flow sensitivity
- Drought curtailment risk

Heat / River Ecology Metrics

- Estimated heat rejection load
- Thermal discharge risk
- River baseline temperature
- Summer temperature stress
- Allowable delta-T threshold
- Dissolved oxygen sensitivity
- Aquatic habitat vulnerability
- Fish intake / entrainment risk where applicable

Energy / Carbon Metrics

- Facility megawatts
- Estimated annual electricity consumption
- Grid carbon intensity
- Renewable energy match percentage
- Cooling energy penalty
- Backup generation risk
- Battery storage / demand response capability

Resilience Metrics

- Seasonal stress multiplier
- Extreme heat vulnerability
- Low-flow percentile
- Grid congestion risk
- Water supply redundancy
- Cooling architecture flexibility

Governance / Permitting Metrics

- Clean Water Act relevance
- NPDES permit relevance
- Cooling water intake relevance
- Thermal discharge review status
- Local water authority constraints
- Missing-data confidence score
- Community concern risk

Sustainability Dashboard Modules

1. Sustainability Snapshot

Shows the overall environmental profile of a proposed facility or river segment.

Outputs:

- Water risk
- Thermal risk
- Energy/carbon risk
- Ecosystem sensitivity
- Regulatory readiness
- Sustainability grade

2. Water Stewardship Panel

Tracks withdrawal, consumption, source type, discharge, and drought exposure.

3. Thermal Ecology Panel

Models heat discharge risk, seasonal river temperature, dissolved oxygen vulnerability, and aquatic habitat sensitivity.

4. Cooling Sustainability Comparator

Compares cooling architectures by water intensity, heat rejection, energy penalty, and climate resilience.

5. Climate-Seasonal Stress Simulator

Shows how the same facility performs under normal, summer peak, drought, and extreme heat scenarios.

6. Sustainability Report Generator

Creates an investor-, regulator-, or public-facing sustainability summary.

MVP Sustainability Score

Create a Sustainability Fit Score from five weighted dimensions:

- Water Stewardship: 25%
- Thermal Ecosystem Safety: 25%
- Energy / Carbon Alignment: 20%
- Climate Resilience: 15%
- Regulatory / Community Readiness: 15%

Sustainability Fit Bands

Score	Classification	Meaning
0–30	Strong Fit	Sustainable buildout candidate
31–55	Conditional Fit	Viable with mitigation

Score	Classification	Meaning
56–75	Sustainability Stress	High environmental/permitting burden
76+	Poor Fit	Redesign, relocate, or defer

Sustainability Agent Actions

Explain Sustainability Tradeoff

Explains whether the proposed buildout is trading water savings for higher electricity demand, or energy savings for higher water consumption.

Recommend Lower-Impact Cooling

Suggests lower-impact cooling options based on river condition, season, and facility load.

Generate Water Stewardship Summary

Creates a formal summary of water source, consumption, discharge, and mitigation strategy.

Generate Climate Resilience Summary

Explains how the project performs during low-flow, summer heat, and drought stress.

Flag Greenwashing Risk

Identifies claims that need stronger evidence, such as vague “sustainable cooling” or “water efficient” statements without supporting metrics.

Draft Public-Facing Sustainability Language

Produces plain-language sustainability explanations for community meetings, public reports, or investor materials.

Sustainability Positioning Statement

RiverCompute Intelligence helps the AI infrastructure market grow without blindly overloading local watersheds. It gives developers, utilities, regulators, and investors a shared evidence base for deciding where compute should go, how it should be cooled, what risks must be mitigated, and when a site should be redesigned before environmental damage or regulatory conflict occurs.

Senior Doctoral Knowledge Package — Economic Sustainability + Agriculture

Expert Positioning

RiverCompute Intelligence should be framed as a Water–Energy–Food–Compute Nexus platform. Its intellectual power is not simply that data centers use water. The deeper argument is that AI infrastructure is becoming a new industrial claimant on the same water, energy, land, and permitting systems that already support agriculture, food security, river ecosystems, utilities, and regional economic development.

Doctoral Thesis

The next era of sustainable economic development will be shaped by whether regions can allocate water, energy, land, and thermal capacity across agriculture, communities, ecosystems, and AI infrastructure without undermining food security, ecological resilience, or long-term regional competitiveness.

RiverCompute Intelligence provides the missing planning layer: a decision intelligence system that evaluates where AI compute can grow without destabilizing watershed health, agricultural resilience, utility capacity, or regulatory trust.

Core Argument

AI data centers should not be evaluated as isolated real estate projects. They should be evaluated as watershed actors, energy actors, agricultural competitors, and regional economic development assets.

That means a proposed hyperscale facility should be tested against five questions:

1. Can the river segment absorb the heat?
2. Can the watershed support the withdrawal and discharge pattern?
3. Will the energy demand alter regional affordability or grid reliability?
4. Could the project compete with agricultural water security during drought or low-flow seasons?
5. Will the project produce enough public value to justify the ecological and infrastructure burden?

Water–Energy–Food–Compute Nexus

Water

Data centers may require water for cooling, humidification, discharge management, and indirectly through electricity production.

Energy

AI data centers can require hundreds of megawatts of continuous power, creating grid, generation, transmission, and affordability implications.

Food / Agriculture

Agriculture remains the dominant global freshwater user and is directly exposed to drought, climate variability, irrigation cost, water rights, and water allocation politics.

Compute

Compute is now a strategic economic input. Regions want AI infrastructure, but compute growth cannot be treated as automatically sustainable.

Agriculture Relevance

Agriculture belongs in the RiverCompute model because it is the benchmark sector for water dependency, land resilience, and food-system economic value. Any serious sustainability model must compare data center water and energy demand against agricultural productivity, irrigation requirements, rural livelihoods, crop value, and drought vulnerability.

RiverCompute should help answer:

- Does a proposed compute facility create water competition with agriculture?
- Does it increase utility pressure during irrigation or heat-stress seasons?
- Can waste heat, reclaimed water, or infrastructure investment create agricultural co-benefits?
- Could data center tax revenue fund watershed resilience, precision irrigation, or rural broadband?
- Is the project extracting regional capacity or building regional resilience?

Economic Sustainability Lens

Economic sustainability is not simply growth. It is growth that does not degrade the underlying resource systems that future economic activity depends on.

For this subject, economic sustainability means:

- AI infrastructure growth without watershed degradation
- Data center investment without agricultural displacement
- Energy expansion without unaffordable rate pressure
- Water allocation without food-system instability
- Industrial development without public trust collapse
- Regional competitiveness without ecological overshoot

Doctoral-Level Model Structure

1. Resource Carrying Capacity

Measures the watershed's ability to support additional compute load.

Inputs:

- River flow
- Low-flow percentile
- Temperature baseline
- Dissolved oxygen
- Ecological sensitivity
- Withdrawal/discharge limits

2. Agricultural Water Security

Measures whether data center demand could stress agricultural production.

Inputs:

- Irrigated acreage
- Crop water demand
- Drought exposure
- Water source competition
- Seasonal irrigation calendar
- Commodity value at risk

3. Energy Burden and Reliability

Measures whether the project strengthens or strains regional energy systems.

Inputs:

- Facility MW
- Local load growth
- Grid capacity
- renewable match
- backup fuel reliance
- transmission constraints
- ratepayer exposure

4. Thermal-Ecological Risk

Measures whether heat rejection could push ecosystems beyond safe thresholds.

Inputs:

- Thermal discharge
- river temperature
- seasonal heat stress
- dissolved oxygen
- aquatic habitat sensitivity
- discharge permit thresholds

5. Regional Net Benefit

Measures whether the project produces enough economic value relative to public resource burden.

Inputs:

- Jobs
- tax revenue
- infrastructure investments
- local procurement
- agricultural co-benefits
- mitigation commitments
- public cost exposure

Signature Framework

WEFC-R Score

Water–Energy–Food–Compute Resilience Score

Weighted dimensions:

- Watershed Capacity: 25%
- Agricultural Water Security: 20%
- Energy Reliability: 20%
- Thermal-Ecological Safety: 20%
- Regional Net Benefit: 15%

WEFC-R Bands

Score	Classification	Meaning
0–30	Resilient Fit	Strong candidate for sustainable buildout
31–55	Conditional Fit	Viable with mitigation and monitoring
56–75	Cross-Sector Stress	Risks to water, agriculture, energy, or ecology require redesign
76+	Unsustainable Allocation	Site should be avoided, delayed, or deeply restructured

High-Impact Research Questions

1. How should river basins allocate thermal capacity between energy, agriculture, communities, and AI infrastructure?
2. Can AI data centers become circular infrastructure assets instead of pure resource consumers?
3. Should data center permitting include agricultural water security impact review?
4. How should states price the opportunity cost of water used for compute versus food production?

5. Can data center waste heat support controlled-environment agriculture, greenhouse systems, aquaculture, or district heating?
6. How can utility planning prevent data center growth from shifting costs onto rural or agricultural communities?
7. What metrics should define a sustainable hyperscale density for a river basin?

Strategic Additions to RiverCompute MVP

Agriculture Impact Module

Adds crop, irrigation, drought, and food-system variables to the model.

Water Allocation Conflict Index

Measures the likelihood that data center water demand competes with agriculture, municipal supply, or ecological flows.

Agricultural Co-Benefit Planner

Identifies opportunities such as reclaimed water use, waste heat reuse, greenhouse partnerships, precision irrigation funding, and rural infrastructure investment.

Regional Net Benefit Calculator

Compares tax revenue, jobs, infrastructure investment, water burden, energy burden, and agricultural risk.

Public Trust and Permit Friction Score

Measures social license risk based on water stress, energy affordability, ecological sensitivity, disclosure quality, and community benefit.

Signature Language

RiverCompute Intelligence does not ask whether a region has enough water in the abstract. It asks whether the watershed, agricultural economy, energy grid, and regulatory system can sustainably absorb the next unit of AI compute.

The future of AI infrastructure will not be decided only by chip supply or land availability. It will be decided by water law, river temperature, agricultural resilience, grid capacity, and public trust.

Closing Position

The regions that win the AI economy will not simply be the regions that build the most data centers. They will be the regions that build compute capacity without sacrificing water security, agricultural productivity, ecological health, or long-term economic resilience.

High-Level Research Package

Working Title

RiverCompute Intelligence: A Water–Energy–Food–Compute Nexus Framework for Sustainable AI Infrastructure Siting

Executive Abstract

Artificial intelligence infrastructure is becoming a new class of regional industrial load. Hyperscale data centers require continuous electricity, cooling capacity, water planning, thermal management, and regulatory approval. Traditional site-selection frameworks often prioritize land, tax incentives, fiber access, and power availability while under-modeling watershed carrying capacity, agricultural water security, thermal discharge, community trust, and ecological resilience.

RiverCompute Intelligence proposes a decision intelligence framework for evaluating whether a river segment, watershed, agricultural economy, and energy system can sustainably absorb additional AI compute. The framework integrates hydrological data, cooling architecture, thermal risk, water withdrawal and consumptive use, seasonal stress, agricultural exposure, regulatory thresholds, and regional net benefit into a practical scoring system for developers, utilities, regulators, investors, and public-sector planners.

Central Research Claim

The next generation of sustainable economic development will depend on whether regions can grow AI infrastructure without degrading the resource systems that make long-term economic activity possible: water, energy, agriculture, ecological health, and public trust.

Core Problem Statement

AI data center expansion is accelerating faster than local planning systems can evaluate cumulative resource pressure. The critical constraint is not simply water volume. It is the interaction between electricity demand, cooling architecture, heat rejection, river temperature, low-flow periods, agricultural water demand, grid capacity, and permitting risk.

Research Gap

Current data center sustainability analysis is fragmented across separate domains:

- Energy consumption and grid load
- Data center cooling efficiency
- Water Usage Effectiveness
- Clean Water Act permitting
- Local economic development

- Agricultural water use
- Climate resilience

The missing layer is a cross-sector carrying-capacity model that evaluates the combined pressure of AI compute on watershed systems, agricultural economies, energy systems, and regional resilience.

Research Objectives

1. Develop a watershed-based compute sustainability framework for hyperscale AI infrastructure.
2. Define measurable indicators for thermal risk, water demand, agricultural exposure, energy burden, and regulatory readiness.
3. Create a Maximum Sustainable Hyperscale Density Index for river segments.
4. Extend the model into a Water–Energy–Food–Compute Resilience Score.
5. Build an MVP decision system capable of producing scenario-level sustainability and regulatory-readiness reports.
6. Demonstrate the framework through a Louisville / Ohio River case study.

Primary Research Questions

1. How much AI compute can a river segment absorb before thermal, water, ecological, or regulatory stress becomes unacceptable?
2. How should data center water demand be evaluated against agricultural water security and food-system resilience?
3. Which cooling architectures minimize total sustainability burden when both direct water use and indirect energy-water tradeoffs are included?
4. How should seasonal low-flow, high-temperature, and drought conditions alter data center siting decisions?
5. Can data centers be designed as circular infrastructure assets through waste heat reuse, reclaimed water, grid services, or agricultural co-benefits?
6. What public value must a hyperscale data center return to justify its resource burden?
7. How should regulators and utilities evaluate cumulative data center clustering along shared river systems?

Key Terms

Watershed-Based Compute Governance

A planning approach that treats data centers as watershed actors whose cooling, water, thermal, and energy impacts must be evaluated within river-basin limits.

Maximum Sustainable Hyperscale Density Index

A score estimating how much data center capacity a river segment can support before ecological, regulatory, or infrastructure stress becomes excessive.

Water–Energy–Food–Compute Nexus

A framework recognizing that AI compute competes and interacts with water systems, power systems, agriculture, food security, and regional economic development.

Thermal Carrying Capacity

The ability of a water body to absorb heat discharge without violating ecological, regulatory, or dissolved oxygen safety conditions.

Agricultural Water Security

The ability of a regional agricultural system to maintain irrigation, crop production, livestock support, and rural economic stability under competing water demands and climate stress.

Conceptual Framework

Layer 1: Watershed Capacity

Measures hydrological ability to support additional industrial load.

Indicators:

- Mean flow
- Low-flow percentile
- river temperature
- seasonal variability
- dissolved oxygen
- drought exposure
- discharge capacity

Layer 2: Cooling Architecture

Measures how facility design affects water, energy, and thermal outcomes.

Indicators:

- Cooling type
- WUE
- energy penalty
- thermal discharge risk
- liquid cooling readiness
- dry/hybrid cooling feasibility

Layer 3: Agricultural Exposure

Measures whether compute expansion creates risk for food and rural economies.

Indicators:

- irrigated acreage
- crop value
- irrigation season overlap
- drought vulnerability
- water source competition
- agricultural employment and revenue exposure

Layer 4: Energy and Grid Burden

Measures whether facility load strengthens or strains regional power systems.

Indicators:

- facility MW
- annual electricity use
- interconnection capacity
- grid congestion
- renewable match
- backup generation
- ratepayer exposure

Layer 5: Regulatory and Public Trust

Measures whether the project is likely to face legal, environmental, or social-license friction.

Indicators:

- Clean Water Act relevance
- NPDES relevance
- withdrawal/discharge thresholds
- aquatic impact exposure
- public disclosure quality
- community benefit commitments

Layer 6: Regional Net Benefit

Measures whether the development returns enough value to justify resource use.

Indicators:

- construction jobs
- permanent jobs
- tax revenue
- infrastructure upgrades
- local procurement
- agriculture co-benefits

- watershed investments

Signature Framework: WEFC-R Score

Water-Energy-Food-Compute Resilience Score

Dimension	Weight	Purpose
Watershed Capacity	25%	Tests hydrological and thermal fit
Agricultural Water Security	20%	Tests food-system and irrigation exposure
Energy Reliability	20%	Tests grid and power-system stress
Thermal-Ecological Safety	20%	Tests heat discharge and ecosystem impact
Regional Net Benefit	15%	Tests economic justification and public value

WEFC-R Classification

Score	Classification	Interpretation
0–30	Resilient Fit	Strong sustainable buildout candidate
31–55	Conditional Fit	Viable with mitigation and monitoring
56–75	Cross-Sector Stress	Requires redesign, restrictions, or stronger public benefit
76+	Unsustainable Allocation	Avoid, defer, or fundamentally restructure the project

Proposed Model Equations

Composite Risk Score

Composite Risk Score =

- Thermal Load Risk $\times 0.35$
- Water Withdrawal Density $\times 0.25$
- Seasonal Overlap Risk $\times 0.20$
- Infrastructure Redundancy Risk $\times 0.20$

Sustainability Fit Score

Sustainability Fit Score =

- Water Stewardship $\times 0.25$
- Thermal Ecosystem Safety $\times 0.25$
- Energy / Carbon Alignment $\times 0.20$
- Climate Resilience $\times 0.15$
- Regulatory / Community Readiness $\times 0.15$

WEFC-R Score

WEFC-R =

- Watershed Capacity $\times 0.25$
- Agricultural Water Security $\times 0.20$
- Energy Reliability $\times 0.20$
- Thermal-Ecological Safety $\times 0.20$
- Regional Net Benefit $\times 0.15$

Literature and Evidence Map

AI Energy Demand

Use energy-demand research to establish urgency and growth pressure.

Data Center Water Use

Use WUE, direct water use, indirect energy-water use, and cooling-system tradeoffs.

Agriculture and Water Security

Use irrigation withdrawal, crop value, and drought vulnerability to connect compute expansion to food-system economics.

Clean Water Act and Thermal Discharge

Use regulatory frameworks around cooling-water intake, thermal discharge, and NPDES review.

Basin-Level Water Governance

Use river basin commissions and watershed authorities as evidence that data center review is moving toward basin-scale regulation.

Louisville / Ohio River Case Study

Use Louisville hyperscale development, Ohio River monitoring, Kentucky agriculture, and ORSANCO water-quality data to ground the model locally.

Proposed Case Study: Louisville / Ohio River

Why Louisville

Louisville is a strong pilot because it connects:

- Active hyperscale data center development
- Ohio River water system

- Regional power infrastructure
- Kentucky agricultural economy
- Public water quality monitoring
- Economic development pressure
- Regulatory and sustainability questions

Case Study Questions

1. What is the resource burden of a 400 MW hyperscale data center campus?
2. How would seasonal Ohio River conditions affect cooling and thermal discharge risk?
3. How does the facility's energy demand compare with regional grid and utility capacity?
4. Could the project create or reduce stress on agricultural water security?
5. What mitigation package would make the development more sustainable?
6. What should a regulator-ready sustainability report include?

Data Architecture

Required Datasets

1. River segment hydrology
2. Water temperature and flow data
3. Data center load and cooling architecture
4. Electricity grid and utility capacity
5. Agricultural land and crop-value data
6. Drought and climate stress data
7. Regulatory thresholds
8. Water withdrawal and discharge permits
9. Community benefit and economic development data

Recommended Data Sources

- USGS Water Data APIs
- ORSANCO Ohio River water quality reports
- EPA Clean Water Act and NPDES materials
- USDA ERS irrigation and farm-sector data
- Kentucky Department of Agriculture economic reports
- Utility announcements and integrated resource planning documents
- Basin commission data center guidance

MVP Research Outputs

1. Technical White Paper

A formal explanation of the Water–Energy–Food–Compute Nexus and the RiverCompute decision framework.

2. Louisville / Ohio River Case Brief

A region-specific pilot showing how the model evaluates a real hyperscale development context.

3. Policy Brief

A concise document for regulators and state agencies explaining why data center permitting should include watershed and agricultural resilience indicators.

4. Investor Risk Memo

A finance-facing memo translating water, energy, and regulatory risks into project risk language.

5. Data Dictionary

A structured glossary of every field needed for the MVP.

6. Scenario Report Template

A reusable report for each proposed data center scenario.

7. Stakeholder Deck

A high-impact slide narrative for agencies, utilities, foundations, data center developers, and sustainability partners.

Research Package Deliverables

Deliverable A: Executive Research Brief

8–12 pages; written for decision-makers.

Deliverable B: Academic-Style Concept Paper

15–25 pages; written for doctoral, grant, or conference positioning.

Deliverable C: MVP Technical Specification

Product-facing build document with datasets, formulas, workflows, and agent actions.

Deliverable D: Policy Adoption Brief

State-facing document proposing RiverCompute metrics as planning criteria.

Deliverable E: Case Study Appendix

Louisville / Ohio River evidence pack.

Grant and Fellowship Framing

Research Category

Sustainable AI infrastructure, water governance, economic sustainability, agricultural resilience, climate adaptation, and regional planning.

Proposal Hook

AI infrastructure is expanding into the same resource systems that sustain food production, water security, and regional energy reliability. This project creates a decision intelligence framework to evaluate whether AI compute can grow without undermining those systems.

Broader Impacts

- Improves sustainable data center planning
- Protects water and agricultural systems
- Supports climate-resilient economic development
- Helps regulators make evidence-based decisions
- Encourages transparent infrastructure governance
- Creates a replicable model for other river basins

High-Level Methodology

Step 1: Literature Review

Review data center energy demand, cooling technologies, water use, agricultural water security, and environmental regulation.

Step 2: Indicator Development

Define measurable variables for watershed capacity, agricultural risk, energy burden, thermal safety, and net benefit.

Step 3: Model Design

Build the MSHDI and WEFC-R scoring systems.

Step 4: Case Study Calibration

Apply the model to Louisville / Ohio River using public datasets and available development information.

Step 5: MVP Implementation

Translate the model into a dashboard, AI agent, scenario builder, and report generator.

Step 6: Validation

Compare model outputs against known regulatory thresholds, real hydrological data, expert review, and pilot stakeholder feedback.

High-Impact Thesis Paragraph

AI infrastructure is becoming a defining test of regional economic sustainability. The same watersheds that support agriculture, ecosystems, cities, and energy systems are now being asked to support large-scale compute. RiverCompute Intelligence reframes the data center siting problem as a Water–Energy–Food–Compute Nexus problem, where the question is not only whether a data center can be built, but whether it can be built without destabilizing the resource systems that make long-term economic resilience possible.

Final Positioning Statement

RiverCompute Intelligence is a watershed-based decision intelligence platform for sustainable AI infrastructure. It helps public and private stakeholders determine where compute should grow, how it should be cooled, what risks must be mitigated, and when a proposed site creates unacceptable pressure on water systems, agriculture, ecological health, energy reliability, or public trust.